

A world map with a heatmap overlay showing rainfall patterns. The colors range from light blue (low rainfall) to dark brown (high rainfall). High rainfall areas are concentrated in the tropics, particularly over the Amazon basin, central Africa, and parts of Southeast Asia. Lower rainfall areas are seen in the subtropics and some high-latitude regions.

Modeling Extreme Events with PyMC

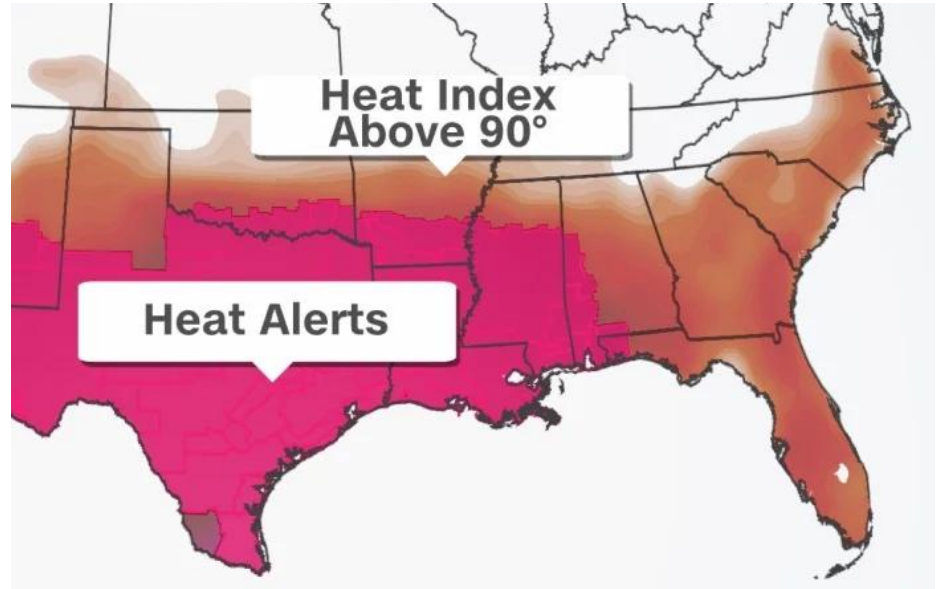
A Bayesian approach to modeling extreme rainfall

Jorn Mossel, Ph.D.
PyData Global 2023

Recent Extreme Weather Events

- **Extreme Temperatures**
- Droughts
- Flash Floods

‘Ridiculous’ heat keeps tormenting Texas, with no end in sight




August 2023

Recent Extreme Weather Events

- Extreme Temperatures
- **Droughts**
- Flash Floods

***How Bad Is the Western Drought?
Worst in 12 Centuries, Study Finds.***



February 2022

Recent Extreme Weather Events

- Extreme Temperatures
- Droughts
- Flash Floods


At least 43 are dead after Ida causes flooding in four states.



September 2021

Overview


1. Statistics of Extreme Events
2. Intro to Bayesian Modelling with PyMC
3. Example: Extreme Rainfall in NYC
4. Bonus example: Gaussian Processes




Statistics of Extreme Events

Introduction to Extreme Value Theory

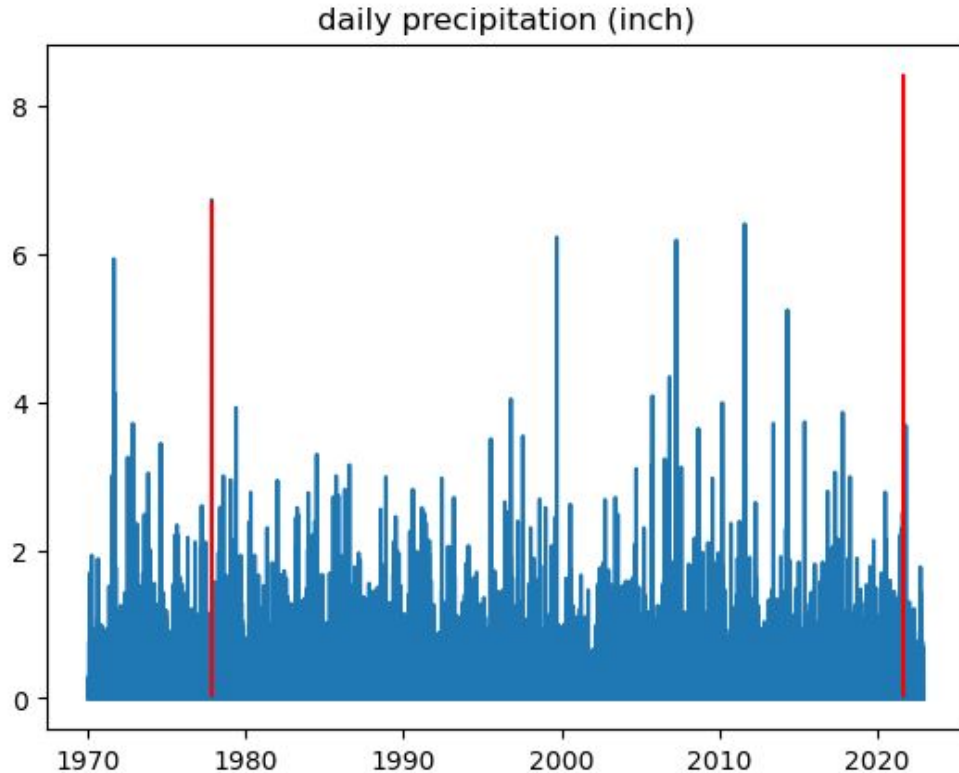
Record rainfall at nearby Newark airport after Hurricane Ida



Record rainfall at nearby Newark airport after Hurricane Ida



How can we define what is extreme?



Q: How extreme is this peak?

A: A better question is to ask how often we expect to observe a level of this or higher?


Q: Isn't that only possible if we have really long history?

A: Not if you know the probability distribution

Q: How do you know which distribution to fit?


A: Extreme Value Theory will tell us

Finding a distribution for the extremes



Complicated (unknown) distribution

Keep only
annual maxima



Extreme Value distribution (see next slide)

Extreme Value Theory states that

*The probability distribution of **maxima** (e.g. annual maxima of daily observations) can be approximated (under general assumptions) with a **Extreme Value distribution** which has a CDF of the form:*


$$G(z) = \exp \left\{ - \left[1 + \xi \left(\frac{z - \mu}{\sigma} \right) \right]^{-1/\xi} \right\}$$

Regardless of what the underlying process is. Which can be max precipitation, temperature extremes, stock market crashes, etc.

Our task is to fit the parameters of this distribution to the observed data

- μ “location”
- σ “scale”
- ξ “shape (tailness)”

Intro to Bayesian Modelling with PyMC



Intro to Bayesian Modelling


Non-bayesian methods

Most ML methods are solving for point estimates of model parameters by

Minimizing a loss function with gradient descent

With this approach the following is hard (in general)

- Uncertainty estimation (eg confidence intervals)
- Taking prior information into account



Intro to Bayesian Modelling

Bayesian methods


Find a *posterior distribution* for the model parameters θ

$$P(\theta|X) = \frac{P(X|\theta)P(\theta)}{\int P(X|\theta)P(\theta)d\theta}$$

Given data X

- *Uncertainty* is captured in $P(\theta|X)$
- *Prior information* can be put into $P(\theta)$

“Fitting” is done via sampling from the distribution



What is PyMC?

PyMC is an open source library for probabilistic programming in Python.

- Intuitive way to express Bayesian models in code
- Fit the model using Monte Carlo Markov Chain (MCMC) sampling
- Includes tools for visualization and diagnostics






Example: Extreme Rainfall in NYC

Fit the Extreme Value Distribution in PyMC

$$P(\mu, \sigma, \xi|z) \propto P(z|\mu, \sigma, \xi)P(\mu, \sigma, \xi)$$

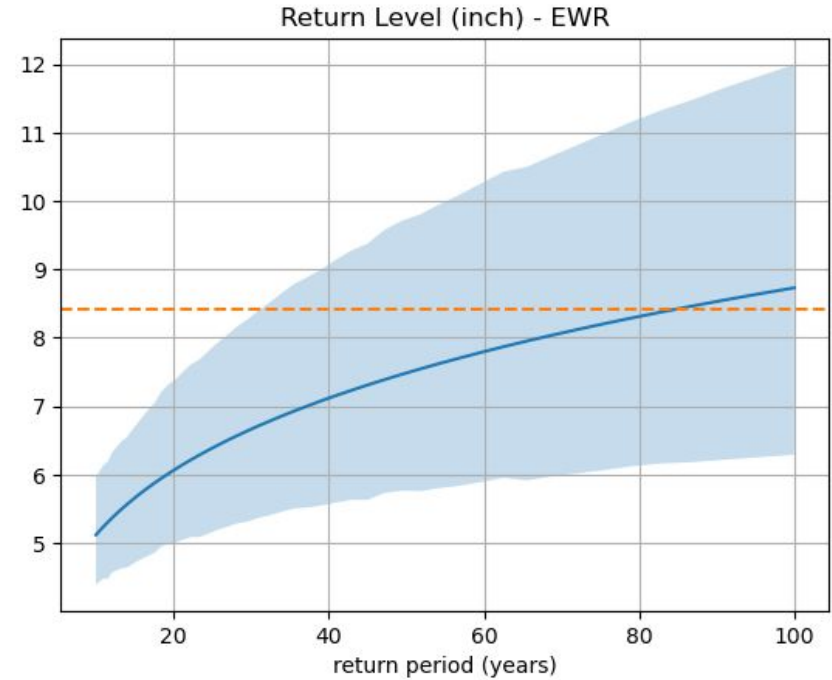
```
with pm.Model() as model:
    # Priors
    μ = pm.Normal("μ", mu=3, sigma=3.0)
    σ = pm.HalfNormal("σ", sigma=1.0)
    ξ = pm.Normal("ξ", mu=0.0, sigma=0.2)
    # Extreme Value likelihood
    gev = pmx.GenExtreme("gev", mu=μ, sigma=σ, xi=ξ, observed=z)
    # MCMC sampling
    trace = pm.sample(2000, target_accept = 0.98)
```

Weakly informative priors are used with the aim to speed up the sampling without biasing the results for the parameters.



Return Levels


- For a **return period** (say every 100 years) we expect to exceed the **return level**



$$P(Z > \text{return level}) = \frac{1}{\text{return period}}$$

Return Levels

- For a **return period** (say every 100 years) we expect to exceed the **return level**
- After fitting the distribution in PyMC it's straightforward to compute statistics like return levels.




```
with model:
```

```
    rp = pm.ConstantData("rp", return_periods)
    rl = pm.Deterministic("rl",  $\mu - \sigma/\xi * (1 - (-np.log(1 - 1/rp)) ** (-\xi))$ )
    posterior_pred = pm.sample_posterior_predictive(trace, var_names=['rl'])
```

Return Levels

- For a **return period** (say every 100 years) we expect to exceed the **return level**
- After fitting the distribution in PyMC it's straightforward to compute statistics like return levels.
- The Newark record (8.4 inch) corresponds to a return period of 83 years.




```
with model:
```

```
    rp = pm.ConstantData("rp", return_periods)
    rl = pm.Deterministic("rl",  $\mu - \sigma/\xi * (1 - (-np.log(1 - 1/rp)) ** (-\xi))$ )
    posterior_pred = pm.sample_posterior_predictive(trace, var_names=['rl'])
```


Return Levels

- For a **return period** (say every 100 years) we expect to exceed the **return level**
- After fitting the distribution in PyMC it's straightforward to compute statistics like return levels.
- The Newark record (8.4 inch) corresponds to a return period of 83 years.
- The advantage of using a Bayesian approach is that we can easily estimate the uncertainty as well

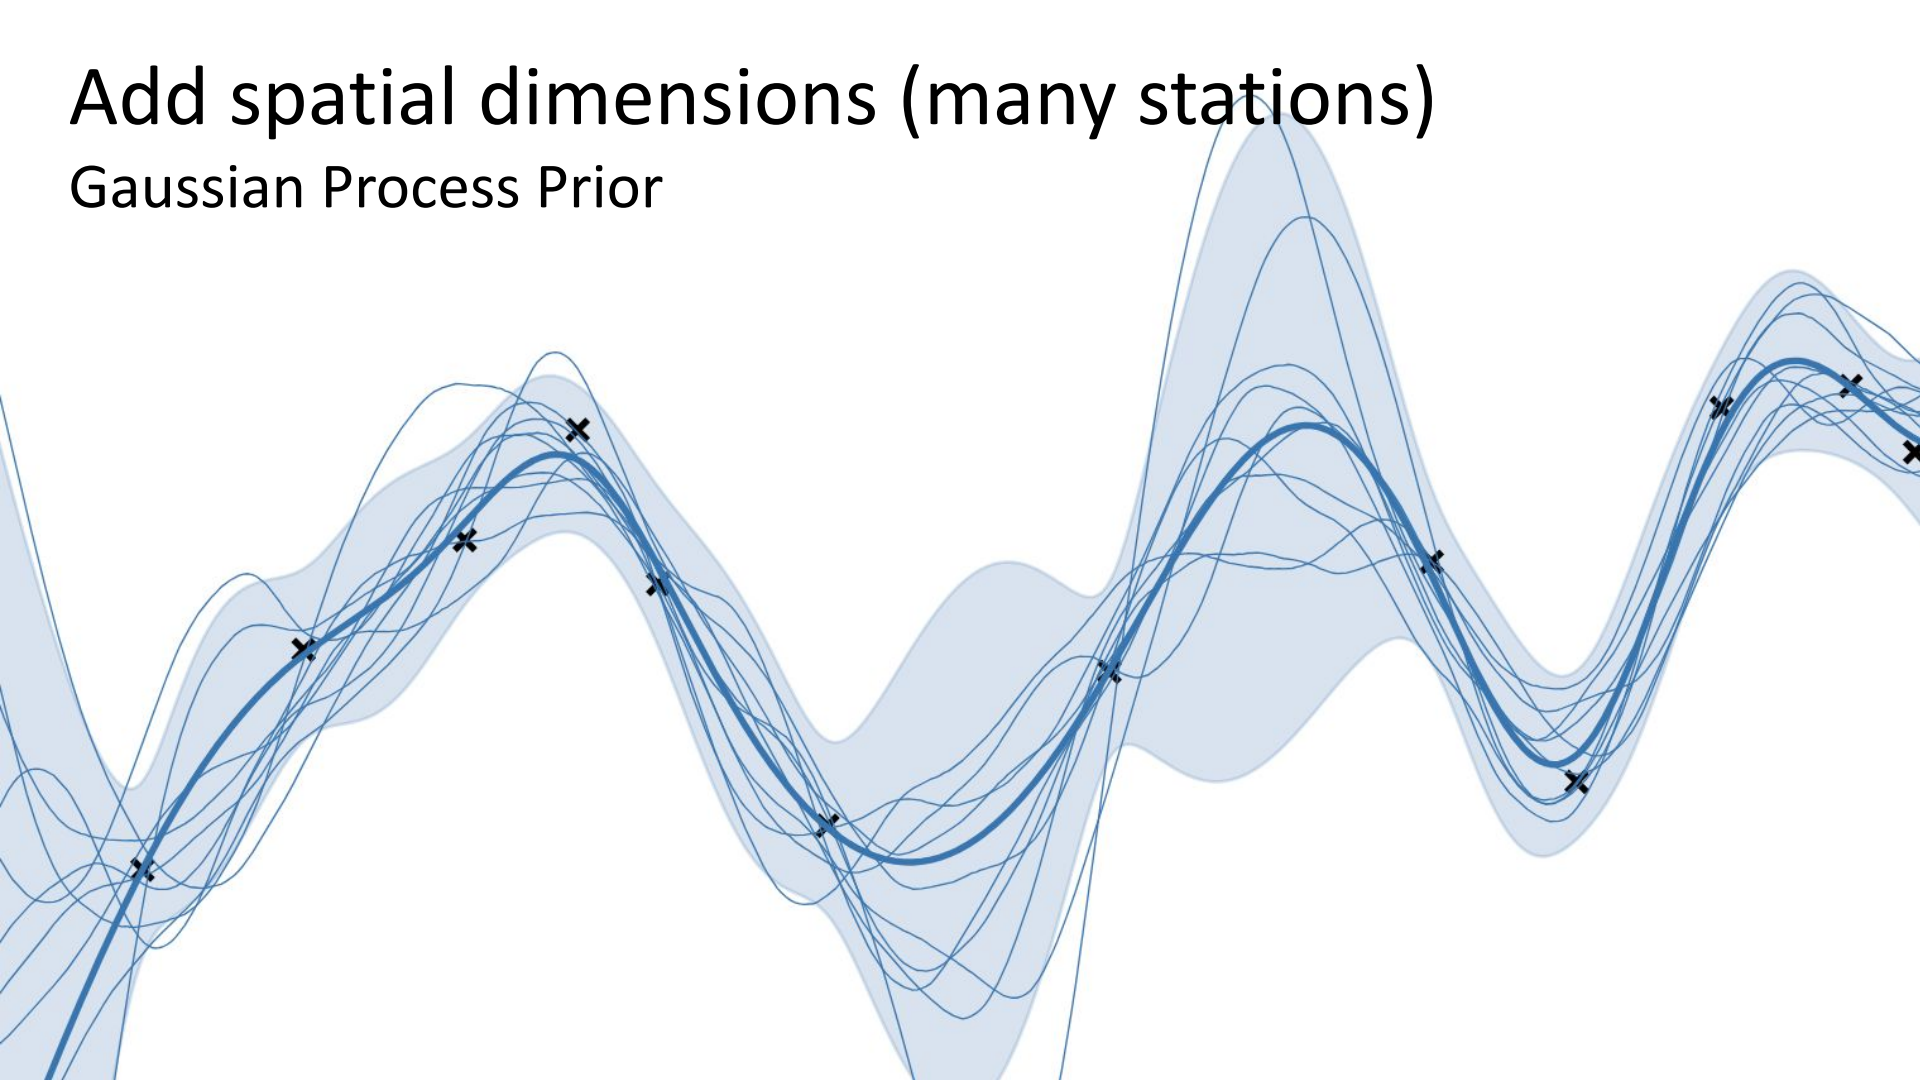


```
with model:
```

```
    rp = pm.ConstantData("rp", return_periods)
    rl = pm.Deterministic("rl",  $\mu - \sigma/\xi * (1 - (-np.log(1 - 1/rp)) ** (-\xi))$ )
    posterior_pred = pm.sample_posterior_predictive(trace, var_names=['rl'])
```

Add spatial dimensions (many stations)

Gaussian Process Prior



How can we improve the model? Using a Gaussian Process!

Add longer history?

- Often not available
- Might not be appropriate if the history is very different


How can we improve the model? Using a Gaussian Process!

Add longer history?

- Often not available
- Might not be appropriate if the history is very different

Add data from other weather stations instead!

- We expect that measurements from nearby stations will be correlated (as function of distance)



How can we improve the model? Using a Gaussian Process!

Add longer history?


- Often not available
- Might not be appropriate if the history is very different

Add data from other weather stations instead!

- We expect that measurements from nearby stations will be correlated (as function of distance)

How to model this in a Bayesian way?

- “All” we need to do is to change the prior distribution
- The prior is a so-called Gaussian Process



Introduction to (Latent) Gaussian Processes


Instead of for each station finding independent prior distributions for the parameters $\mu_1, \mu_2, \mu_3, \dots$

Find a prior distribution for a function $\mu(x)$ which can be evaluated at different locations

We want

- The mean of $\mu(x)$ to be smooth
- The variance of $\mu(x)$ to reduce when we add observations

A Gaussian Process achieves exactly this!



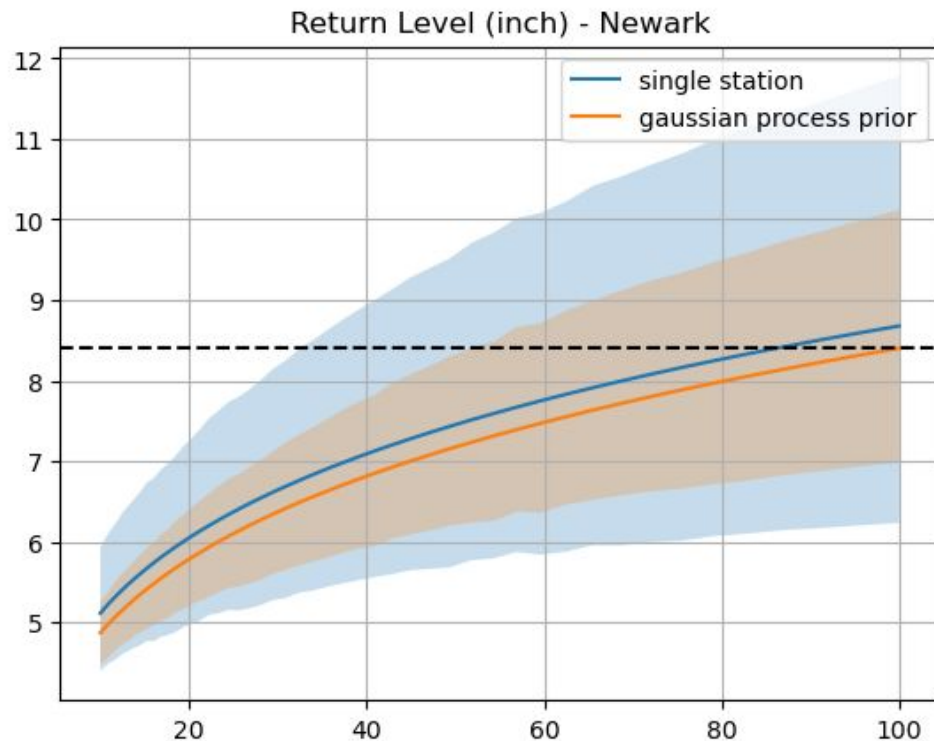
Bayesian Model with Gaussian Process Prior

```
with pm.Model(coords=coords) as gp_model:
    pt_idx = pm.ConstantData("station_idx", station_idx, dims="obs")
    pt_X_lonlat = pm.ConstantData("station_loc", X_lonlat, dims=("station", "lonlat"))
    # gaussian process hyper parameters
     $\iota$  = pm.InverseGamma(" $\iota$ ", mu = 50.0, sigma = 50.0)
     $\eta$  = pm.Gamma(" $\eta$ ", mu=0.15, sigma=0.10, dims = "cov_params")
    # gaussian process prior for  $\mu$ 
    gp_ $\mu$  = pm.gp.Latent(cov_func= $\eta[0]**2$  * Matern32Chordal(2,  $\iota$ ))
     $\mu_{\text{group}}$  = pm.Normal(" $\mu_{\text{group}}$ ", mu=3.0, sigma=1.0)
     $\mu$  = pm.Deterministic(" $\mu$ ",  $\mu_{\text{group}}$  + gp_ $\mu$ .prior(" $\mu_{\text{gp}}$ ", X=pt_X_lonlat), dims="station")
    # gaussian process prior for  $\sigma$ 
    gp_ $\sigma_{\text{log}}$  = pm.gp.Latent(cov_func= $\eta[1]**2$  * Matern32Chordal(2,  $\iota$ ))
     $\sigma_{\text{log}}$  = gp_ $\sigma_{\text{log}}$ .prior(" $\sigma_{\text{log}}$ ", X=pt_X_lonlat, dims="station")
     $\sigma_{\text{log\_group}}$  = pm.Normal(" $\sigma_{\text{log\_group}}$ ", mu=-1.0, sigma=2.0)
     $\sigma$  = pm.Deterministic(" $\sigma$ ", pm.math.exp( $\sigma_{\text{log\_group}}$  +  $\sigma_{\text{log}}$ ), dims="station")
    # gaussian process prior for  $\xi$ 
    gp_ $\xi$  = pm.gp.Latent(cov_func= $\eta[2]**2$  * Matern32Chordal(2,  $\iota$ ))
     $\xi_{\text{group}}$  = pm.TruncatedNormal(" $\xi_{\text{group}}$ ", mu=0.0, sigma=0.25, lower=-0.99, upper=0.99)
     $\xi$  = pm.Deterministic(" $\xi$ ", pm.math.tanh( $\xi_{\text{group}}$  + gp_ $\xi$ .prior(" $\xi_{\text{gp}}$ ", X=pt_X_lonlat)), dims="station")
    # likelihood for all observations
    gev = pmx.GenExtreme("gev", mu= $\mu$ [pt_idx], sigma= $\sigma$ [pt_idx], xi= $\xi$ [pt_idx], observed=y_, dims="obs")
```

The model is “slightly” more complicated but still “fits” on the screen.

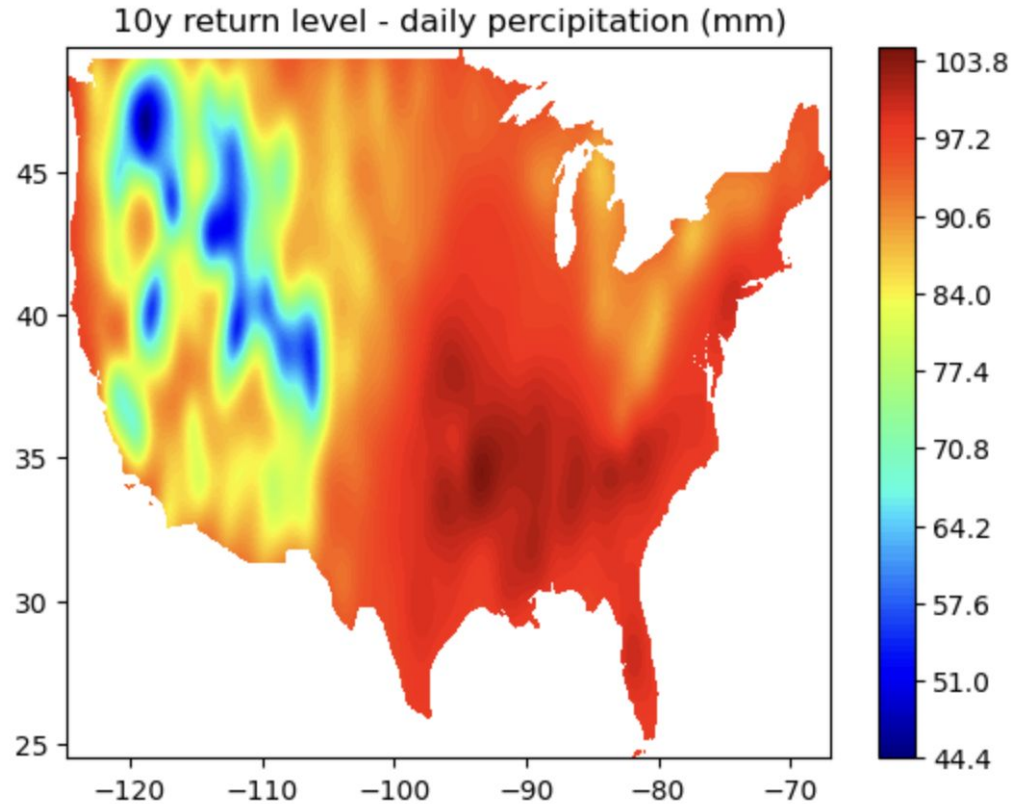
Results for Gaussian Process Prior

- Using 10 nearby stations to get a more accurate prediction for Newark
- Uncertainty almost twice as small!
- New estimate: Hurricane Ida was a once in a **100-year event**



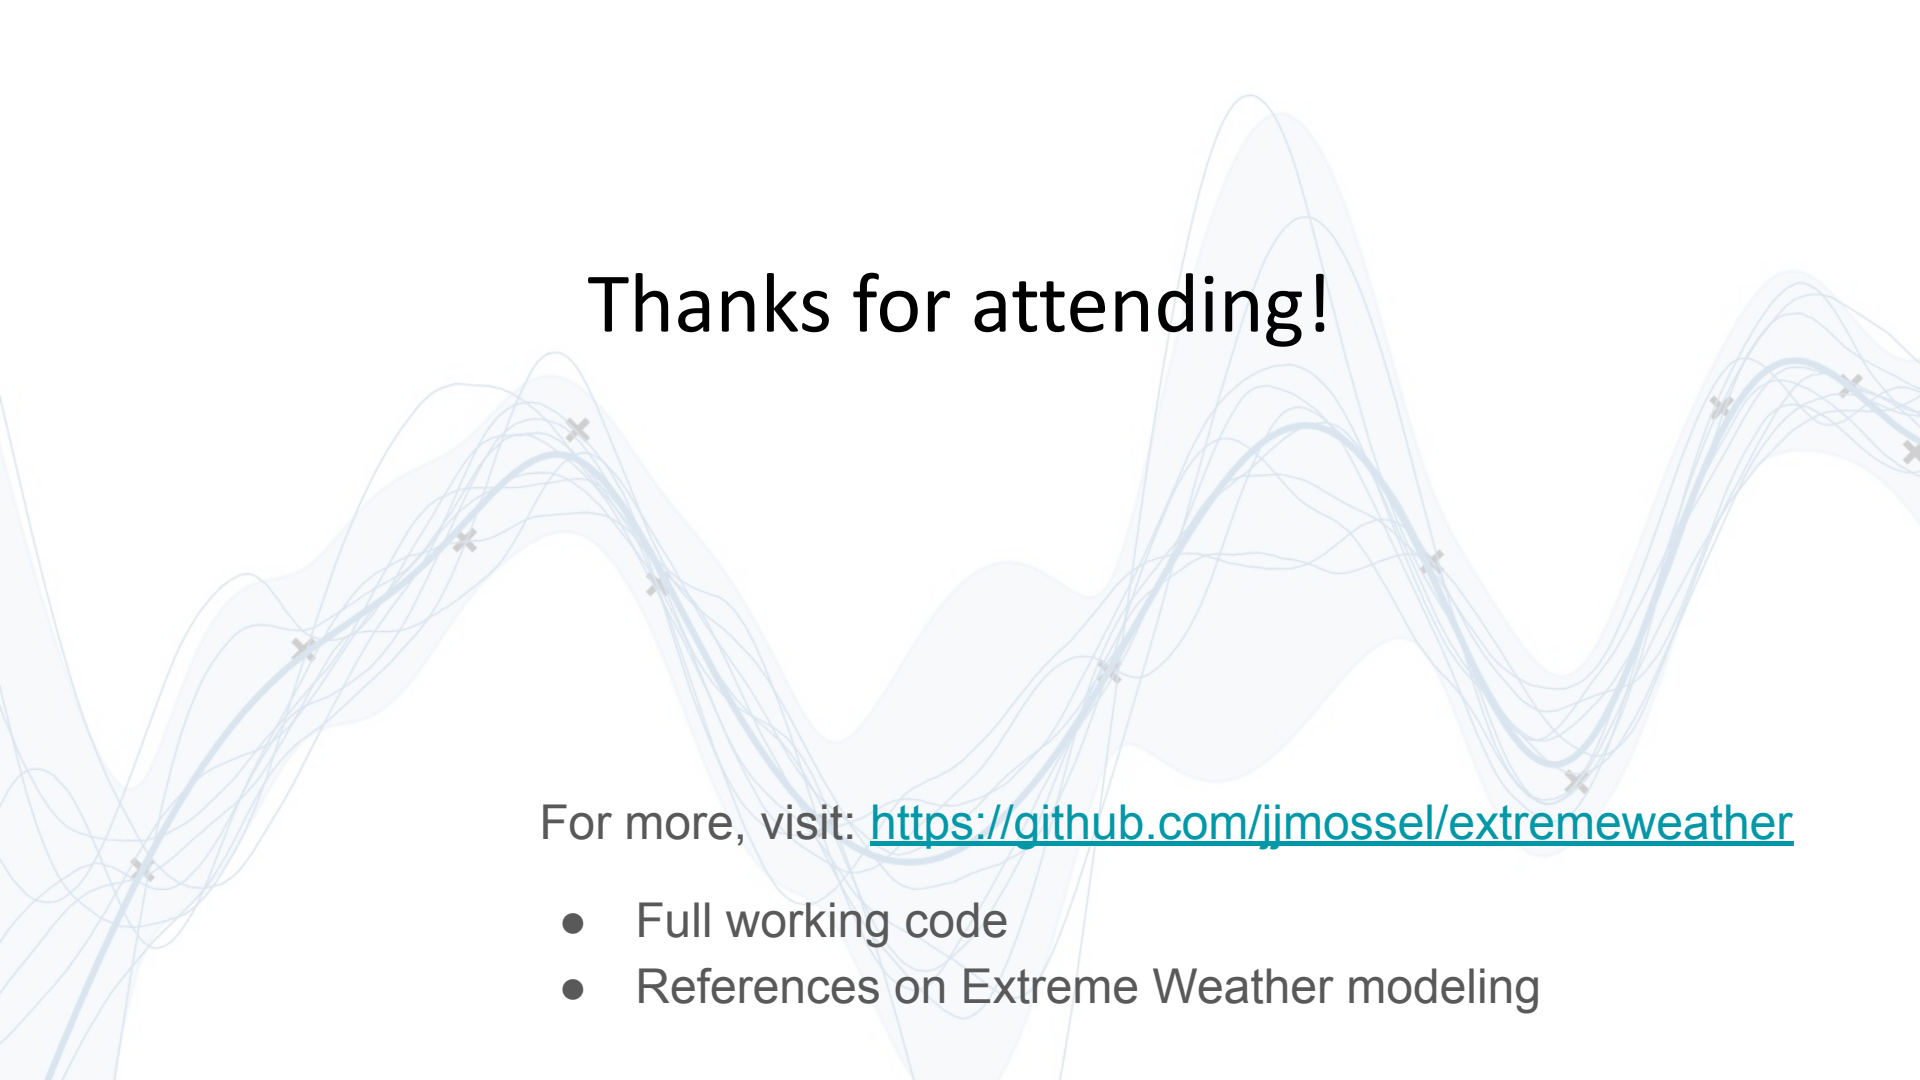
Even more stations

3500 stations



Takeaways

- Modelling Extreme Events -> Extreme Value Theory
- Bayesian Modelling in python -> PyMC
- Modelling spatial problems -> Gaussian Processes



Thanks for attending!

For more, visit: <https://github.com/jjmossel/extremeweather>

- Full working code
- References on Extreme Weather modeling